

YIHAN DU

yihan_du@sutd.edu.sg ◊ <https://yihandu.github.io>

RESEARCH INTERESTS

Machine learning, with emphases on reinforcement learning, online learning (in particular, multi-armed bandit) and representation learning.

EMPLOYMENT

Singapore University of Technology and Design

August 2025 - Present

Pillar of Engineering Systems and Design

Tenure-track Assistant Professor

University of Illinois at Urbana-Champaign

August 2023 - July 2025

Department of Electrical and Computer Engineering

Postdoctoral Researcher

Advisor: Prof. R. Srikant (IEEE Fellow)

EDUCATION

Tsinghua University

September 2018 - June 2023

Institute for Interdisciplinary Information Sciences (headed by Prof. Andrew Chi-Chih Yao)

Ph.D. in Computer Science

Advisor: Prof. Longbo Huang

Dissertation: Risk-aware and Efficient Online Decision Making (**Tsinghua Outstanding Doctoral Dissertation Award**)

Xiamen University

September 2014 - June 2018

B.E. in Computer Science

Rank: 3/93

VISIT & INTERNSHIP

Cornell University

September - December 2022

Visiting scholar (in-person)

Research topic: reinforcement learning

Supervisor: Prof. Wen Sun

Microsoft Research Asia

January - May 2020

Research intern

Research topic: online learning

Supervisor: Dr. Wei Chen (ACM/IEEE Fellow, Director of MSR Asia Theory Center)

PUBLICATIONS

Yihan Du, Anna Winnicki, Gal Dalal, Shie Mannor, R. Srikant, “Reinforcement Learning with Segment Feedback,” International Conference on Machine Learning (ICML), 2025.

Yihan Du, Anna Winnicki, Gal Dalal, Shie Mannor, R. Srikant, “Exploration-Driven Policy Optimization in RLHF: Theoretical Insights on Efficient Data Utilization,” International Conference on Machine Learning (ICML), 2024.

Yihan Du, R. Srikant[†], Wei Chen[†], “Cascading Reinforcement Learning,” International Conference on Learning Representations (ICLR), 2024 ([†]equal advising, **spotlight**).

Yu Chen[#], **Yihan Du**, Pihe Hu, Siwei Wang, Desheng Wu, Longbo Huang, “Provably Efficient Iterated CVaR Reinforcement Learning with Function Approximation and Human Feedback,” International Conference on Learning Representations (ICLR), 2024 ([#]graduate student mentored with my Ph.D. advisor).

Nuoya Xiong[#], **Yihan Du**, Longbo Huang, “Provably Safe Reinforcement Learning with Step-wise Violation Constraints,” Proceedings of the Conference on Neural Information Processing Systems (NeurIPS), 2023 ([#]undergraduate student mentored with my Ph.D. advisor).

Yihan Du, Longbo Huang, Wen Sun, “Multi-task Representation Learning for Pure Exploration in Linear Bandits,” International Conference on Machine Learning (ICML), 2023.

Yihan Du, Siwei Wang, Longbo Huang, “Provably Efficient Risk-Sensitive Reinforcement Learning: Iterated CVaR and Worst Path,” International Conference on Learning Representations (ICLR), 2023.

Yihan Du, Wei Chen, Yuko Kuroki, Longbo Huang, “Collaborative Pure Exploration in Kernel Bandit,” International Conference on Learning Representations (ICLR), 2023.

Yihan Du, Wei Chen, “Branching Reinforcement Learning,” International Conference on Machine Learning (ICML), 2022.

Yihan Du, Siwei Wang, Zhixuan Fang, Longbo Huang, “Continuous Mean-Covariance Bandits,” Proceedings of the Conference on Neural Information Processing Systems (NeurIPS), 2021.

Yihan Du, Yuko Kuroki, Wei Chen, “Combinatorial Pure Exploration with Bottleneck Reward Function,” Proceedings of the Conference on Neural Information Processing Systems (NeurIPS), 2021.

Yihan Du, Siwei Wang, Longbo Huang, “A One-Size-Fits-All Solution to Conservative Bandit Problems,” Proceedings of the AAAI Conference on Artificial Intelligence (AAAI), 2021.

Yihan Du*, Yuko Kuroki*, Wei Chen, “Combinatorial Pure Exploration with Full-Bandit or Partial Linear Feedback,” Proceedings of the AAAI Conference on Artificial Intelligence (AAAI), 2021 (*equal contribution).

(* α - β ordering) Wei Chen, **Yihan Du**, Longbo Huang, Haoyu Zhao, “Combinatorial Pure Exploration for Dueling Bandit,” International Conference on Machine Learning (ICML), 2020.

Yihan Du, Siwei Wang, Longbo Huang, “Dueling Bandits: From Two-dueling to Multi-dueling,” Proceedings of the International Conference on Autonomous Agents and Multiagent Systems (AAMAS), 2020.

Yihan Du, Yan Yan, Si Chen, Yang Hua, “Object-adaptive LSTM Network for Real-time Visual Tracking with Adversarial Data Augmentation,” Neurocomputing, 2019.

Yihan Du, Yan Yan, Si Chen, Yang Hua, Hanzi Wang, “Object-adaptive LSTM Network for Visual Tracking,” International Conference on Pattern Recognition (ICPR), 2018.

SELECTED AWARDS

China Computer Federation (CCF) Agent and Multi-Agent System Doctoral Dissertation Award, by CCF Multi-Agent System Committee, June 2024 (the only recipient nationwide) *June 2024*

Tsinghua Outstanding Doctoral Dissertation Award, by Tsinghua University (the only recipient among CS graduates at IIIS, Tsinghua University in 2023) *June 2023*

Beijing Outstanding Graduate, by Beijing Municipal Education Commission (the only recipient among CS graduates at IIIS, Tsinghua University in 2023) *June 2023*

China National Scholarship for Ph.D. Students, by Ministry of Education of China (the only recipient among CS students at IIIS, Tsinghua University in 2022) *October 2022*

Toyota Scholarship, by Toyota and Tsinghua University

October 2021

Outstanding Graduate, by Xiamen University

June 2018

INVITED TALKS

“Why is RLHF Data-Efficient in Policy Optimization”

- UC Riverside CS

February 2024

- Nanyang Technological University CCDS

January 2024

- Colorado School of Mines CS

January 2024

- New Jersey Institute of Technology CS

December 2024

- China Computer Federation (CCF) Agent and Multi-Agent System Seminar

June 2024

“Risk-aware Online Decision Making”

- TrustML Young Scientist Seminar, RIKEN AIP

May 2023

- MLOPT Idea Seminar, UW-Madison

April 2023

“Combinatorial Pure Exploration for Dueling Bandit”

- CCF Doctoral Forum in Theoretical Computer Science

June 2021

ACADEMIC SERVICE & ACTIVITIES

Reviewer

Conference: ICML, NeurIPS, ICLR, AAAI, UAI, RLC

Journal: Transactions on Pattern Analysis and Machine Intelligence (TPAMI), Journal of Machine Learning Research (JMLR), Transactions on Networking (ToN), Transactions on Machine Learning Research (TMLR), Transactions on Network Science and Engineering (TNSE)

Technical Program Committee (TPC) Member

INFOCOM, IEEE/IFIP WiOpt

Teaching Assistant

Stochastic Network Optimization, graduate course (IIIS, Tsinghua University)

Spring 2021

Introduction to Computer Science, undergraduate course (Yao Class, Tsinghua University)

Fall 2019